

# Recommendation Systems

Cam Tu Nguyen

阮锦绣

Software Institute, Nanjing University  
nguyenct@lamda.nju.edu.cn  
ncamt@gmail.com


# Outline

- Recommendation Systems
- Main Approaches
  - Content-Based Recommendations
  - Collaborative Filtering
- Approximate Nearest Neighbor Search
  - Locality Sensitive Hashing

# Recommendations

- **Product Recommendations**
  - Online retailers such as Amazon, Alibaba
  - Return users products that they might like to buy
- **Movie Recommendations**
  - Netflix offers its customers recommendations of movies they might like.
  - The recommendations are based on ratings provided by users
- **News Articles**
  - News services have attempted to identify articles of interest to readers ,based on the articles that they have read in the past.
- Other applications: blogs recommendations, video recommendations on Youtube, etc.

# The Long Tail



Source: Chris Anderson (2004)

Sources: Erik Brynjolfsson and Jeffrey Hu, MIT, and Michael Smith, Carnegie Mellon; Barnes & Noble; Netflix; RealNetworks

# Recommendation Types

- Editorial
- Simple aggregates
  - Top 10, Most Popular, Recent Uploads
- Tailored to individual users
  - Amazon, Netflix, ...

# Formal Model

- $C$  = set of Customers
- $S$  = set of Items
- Utility function  $u: C \times S \rightarrow R$ 
  - $R$  = set of ratings
  - $R$  is a totally ordered set
  - E.g., 0-5 stars, real number in  $[0,1]$

# Utility Matrix

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

# Populating the Utility Matrix

- **Explicit**
  - Ask people to rate items
  - Doesn't work well in practice – people can't be bothered.
- **Implicit**
  - Learn ratings from user actions
  - E.g., purchase implies high ratings
  - What is about



# Two basic architectures for a recommendation system

- Content-based systems focus on properties of items
  - Similarity of items is determined by measuring the similarity in their properties.
- Collaborative-Filtering systems focus on the relationship between users and items.
  - Similarity of items is determined by the similarity of the ratings of those items by the users who have rated both items.

# Outline

- Recommendation Systems
- Main Approaches
  - Content-Based Recommendations
  - Collaborative Filtering
- Approximate Nearest Neighbor Search
  - Locality Sensitive Hashing

# Item Profiles

- For each item, create an item profile
- Profile is a set of features/attributes
  - Movies: author, title, actor, director, ...
  - Text: set of “important” words in document
  - Music Product: artist, composer, and genre.

# TF.IDF

- How to pick important words?
  - Use heuristic is TF.IDF (Term Frequency times Inverse Doc Frequency)

- TF.IDF

- $f_{ij}$  = frequency of term  $t_i$  in document  $d_j$
- $n_i$  = number of docs that mention term  $t_i$
- $N$  = total number of docs

$$TF_{ij} = \frac{f_{ij}}{\max_k f_{kj}}$$

$$IDF_i = \log \frac{N}{n_i}$$

- TF.IDF score  $w_{ij} = TF_{ij} \times IDF_i$
- Doc profile = set of words with highest TF.IDF scores, together with their scores.

# Representing Item Profiles

- **Example:** Suppose the only features of movies are the set of actors and the average rating. Consider two movies with five actors each.
  - Two of the actors are in both movies
  - One movie has an average rating of 3, and the other an average of 4.

0	1	1	0	1	1	0	1	$3\alpha$
1	1	0	1	0	1	1	0	$4\alpha$

- **Alpha** is introduced as a scaling factor for the average rating feature.
- Similarity between 2 items can be measured using **cosine**

# User profiles and predictions

- User profiles describe users' preferences
- User profile possibilities
  - Average of rated item profiles.
  - Variation: normalize the utilities by subtracting the average value for a user. That way, we get negative weights for items with a below-average rating, and positive weights for items with above-average rating.
  - ...

# Recommending Items to Users Based on Content

- Given user profile  $\mathbf{c}$  and item profile  $\mathbf{s}$ 
  - Estimate  $u(\mathbf{c}, \mathbf{s}) = \cos(\mathbf{c}, \mathbf{s}) = \mathbf{c} \cdot \mathbf{s} / (|\mathbf{c}| |\mathbf{s}|)$
  - Need efficient method to find items with high utility
    - Locality Sensitive Hashing (stay tune!)

# Limitations of content-based approach

- Find the appropriate features
  - E.g., images, movies, music
- Overspecialization
  - Never recommends items outside user's content profile
  - People might have multiple interests
- Recommendations for new users
  - How to build a profile?



# Collaborative Filtering

- Consider user  $c$
- Find set  $D$  of other users whose ratings are “similar” to  $c$ ’s ratings
- Estimate user’s ratings based on ratings of users in  $D$

# Similar Users

- Let  $r_x$  be the vector of user  $x$ 's ratings
- Cosine similarity measure
  - $\text{Sim}(x, y) = \cos(r_x, r_y)$
- Pearson correlation coefficient
  - $S_{xy}$  = items rated by both users  $x$  and  $y$

$$\text{sim}(x, y) = \frac{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)(r_{ys} - \bar{r}_y)}{\sqrt{\sum_{s \in S_{xy}} (r_{xs} - \bar{r}_x)^2 (r_{ys} - \bar{r}_y)^2}}$$

- Other similarity measures: Jaccard Distance, etc.

# Similar Users

	HP1	HP2	HP3	TW	SW1	SW2	SW3
<i>A</i>	4			5	1		
<i>B</i>	5	5	4				
<i>C</i>				2	4	5	
<i>D</i>		3					3

- Cosine Similarity**

- We can treat blanks as a 0 value (it might be not the best choice)
- The cosine of the angle between A and B is

$$\frac{4 \times 5}{\sqrt{4^2 + 5^2 + 1^2} \sqrt{5^2 + 5^2 + 4^2}} = 0.380$$

- The cosine of the angle between A and C is

$$\frac{5 \times 2 + 1 \times 4}{\sqrt{4^2 + 5^2 + 1^2} \sqrt{2^2 + 4^2 + 5^2}} = 0.322$$

# Rating predictions

- Let **D** be the set of **k** users most similar to **c** who have rated item **s**
- Possibilities for prediction function (item **s**):

$$r_{cs} = 1/k \sum_{d \text{ in } D} r_{ds}$$

$$r_{cs} = (\sum_{d \text{ in } D} \text{sim}(c,d) r_{ds}) / (\sum_{d \text{ in } D} \text{sim}(c,d))$$

# Complexity

- Expensive step is finding  $k$  most similar customers
  - For each user,  $O(|U|)$
  - How to make it faster? (Again, Locality Sensitive Hashing comes to rescue!)
- Too expensive to do at runtime
  - Could pre-compute (e.g. using MapReduce in offline mode)
- Can use clustering, partitioning as alternatives, but quality degrades.

# Item-Item Collaborative Filtering

- So far: User-user collaborative filtering
- Another view
  - For item  $s$ , find other similar items
  - Estimate rating for item based on ratings for similar items
  - Can use same similarity metrics and prediction functions as in user-user model.
- In practice, it has been observed that item-item often works better than user-user.

# Pros and cons of Collaborative Filtering

- Works for any kind of item
  - No feature selection needed
- New user problem
- New item problem
- Sparsity of rating matrix
  - Cluster-based smoothing
  - Add more data.

# Evaluating Predictions

- Compare predictions with known ratings
  - Root-mean-square error (RMSE)
- Another approach: 0/1 model
  - Coverage
    - Number of items/users for which system can make predictions
  - Precision
    - Accuracy of predictions
  - Receiver operating characteristic (ROC)
    - Tradeoff curve between false positives and false negatives.



# Finding similar vectors

- Common problem that comes up in many settings
- Given a large number  $N$  of vectors in some high-dimensional space ( $M$  dimensions), find pairs of vectors that have high similarity
  - E.g. User profiles, item profiles
- We need a method to solve this problem fast! (Locality Sensitive Hashing)


# Outline

- Recommendation Systems
- Main Approaches
  - Content-Based Recommendations
  - Collaborative Filtering
- Near Neighbor Search in High Dimensional Data
  - Locality Sensitive Hashing

# Locality Sensitive Functions

- *Locality-sensitive (LS) family is a family of functions that can be combined to distinguish strongly between pairs at a low distance from pairs at a high distance.*
- **Three conditions for a LS family function**
  - They must be more likely to make close pairs be candidate pairs than distant pairs
  - They must be statistically independent
  - They must be efficient, in two ways
    - Serve to identify candidates pairs in time much less than the time it takes to look at all pairs.
    - They must be combinable to build functions that are better at avoiding false positives and negatives.

# A $(d_1, d_2, p_1, p_2)$ -sensitive function



# Amplifying a LS-family

- Two constructions:
  - AND construction
  - OR construction
- **AND of Hash functions**
  - Given family  $H$ , construct family  $H'$  consisting of  $r$  functions from  $H$
  - For  $h=[h_1, \dots, h_r]$  in  $H'$ ,  $h(x)=h(y)$  if and only if  $h_i(x)=h_i(y)$  for all  $i$ .
  - Theorem: If  $H$  is  $(d_1, d_2, p_1, p_2)$ -sensitive, then  $H'$  is  $(d_1, d_2, (p_1)^r, (p_2)^r)$ -sensitive.
- **OR of Hash functions**
  - Given family  $H$ , construct family  $H'$  consisting of  $b$  functions from  $H$
  - For  $h=[h_1, \dots, h_b]$  in  $H'$ ,  $h(x)=h(y)$  if and only if  $h_i(x)=h_i(y)$  for **some**  $i$ .
  - Theorem: If  $H$  is  $(d_1, d_2, p_1, p_2)$ -sensitive, then  $H'$  is  $(d_1, d_2, 1-(1-p_1)^b, 1-(1-p_2)^b)$ -sensitive.

# AND-OR Composition

- Apply a  $r$ -way AND construction followed by an  $b$ -way OR construction.
- Transforms probability  $p$  into  $1-(1-p^r)^b$ .
- **Example:** Take  $H$  and construct  $H'$  by the AND construction with  $r=4$ . Then from  $H'$ , construct  $H''$  by the OR construction with  $b=4$ .

# AND-OR Composition

- **Example:** Take  $H$  and construct  $H'$  by the AND construction with  $r=4$ . Then from  $H'$ , construct  $H''$  by the OR construction with  $b=4$ .

$p$	$1-(1-p^4)^4$
.2	.0064
.3	.0320
.4	.0985
.5	.2275
.6	.4260
.7	.6666
.8	.8785
.9	.9860

**Example:** Transforms a  $(.2,.8,.8,.2)$ -sensitive family into a  $(.2,.8,.8785,.0064)$ -sensitive family.

# OR-AND Composition

- Apply a  $b$ -way OR construction followed by an  $r$ -way AND construction.
- Transforms probability  $p$  into  $(1-(1-p)^b)^r$ .
- **Example:** Take  $H$  and construct  $H'$  by the OR construction with  $b=4$ . Then from  $H'$ , construct  $H''$  by the AND construction with  $r=4$ .



# OR-AND Composition

- **Example:** Take  $H$  and construct  $H'$  by the OR construction with  $b=4$ . Then from  $H'$ , construct  $H''$  by the AND construction with  $r=4$ .

$p$	$(1-(1-p)^4)^4$
.1	.0140
.2	.1215
.3	.3334
.4	.5740
.5	.7725
.6	.9015
.7	.9680
.8	.9936

**Example:** Transforms a  $(.2, .8, .8, .2)$ -sensitive family into a  $(.2, .8, .9936, .1215)$ -sensitive family.

# Summary of LS families

- Pick any two distances  $x < y$
- Start with a  $(x, y, (1-x), (1-y))$ -sensitive family
- Apply constructions to produce  $(x, y, p, q)$ -sensitive family, where  $p$  is almost 1 and  $q$  is almost 0.
- The closer to 0 and 1 we get, the more hash functions must be used.


# LSH for Cosine Distance

- Random Hyper planes
  - Convert data matrix to signature matrix (each signature corresponds to one object).
  - A  $(d_1, d_2, (1-d_1/180), (1-d_2/180))$ -sensitive family for any  $d_1, d_2$ .
- Apply hashing to the signature matrix
  - Objects in the same buckets are candidate pairs.

# Random Hyperplanes

- Pick a random vector  $v$ , which determines a hash function  $h_v$  with two buckets
  - $h_v(x) = +1$  if  $v \cdot x > 0$ ;  $-1$  if  $v \cdot x < 0$
- LS-family  $H$  = set of all functions derived from any vector.
- Claim: For points  $x$  and  $y$ 
  - $P[h(x)=h(y)] = 1 - d(x,y)/180$

# Proof of Claim



# Locality Sensitive Hashing for Cosine Distance


- Signatures for Cosine Distance
  - Pick some number of random vectors, and hash your data for each vector.
    - It suffices to consider only vectors consisting of +1 and -1 components
  - The result is a signature (**sketch**) of +1 and -1's for each data point.

# Locality Sensitive Hashing for Cosine Distance

- **Example:** Suppose our space is 4-dimensional space
  - Consider two vectors  $x=[3,4,5,6]$  and  $y=[4,3,2,1]$
  - The cosine of the angle between  $x$  and  $y$  is 0.7875, or the angle between  $x$  and  $y$  is about 38 degrees.
  - We pick 3 random vectors:  $v1=[+1, -1, +1, +1]$ ;  $v2=[-1, +1, -1, +1]$ , and  $v3=[+1, +1, -1, -1]$ .
    - For the vector  $x=[3,4,5,6]$ , the sketch is  $[+1,+1, -1]$
    - For the vector  $y=[4,3,2,1]$ , the sketch is  $[+1, -1, +1]$
    - The sketches for  $x$  and  $y$  agree in 1/3 of the positions, we estimate the angle between them is 120 degrees. (**not even close!**)
  - If we look at all 16 random vectors (why 16?)
    - There are only 4 of  $v$  vectors where  $v.x$  and  $v.y$  have different signs.
      - The vectors include  $v2$ , and  $v3$
    - The estimate of the angle would have been  $180/4=45$  degrees. (**better**)


# Locality Sensitive Hashing for Cosine Distance

- Signatures for Cosine Distance
  - Pick some number of random vectors, and hash your data for each vector.
  - The result is a signature (**sketch**) of +1 and -1's for each data point.
- LSH: Partition into bands



Signature Matrix:  
#rows = # random vectors  
#columns = # data objects





# Partition into Bands (cont.)

- Divide signature matrix **M** into **b** bands of **r** rows.
  - Create one hash table per band.
- For each band, hash its portion of each column to its hash table
- ***Candidate pairs*** are columns that hash to the same bucket for  $\geq$  band.
- Tune **b** and **r** to catch most similar pairs, but few non-similar pairs.

# Other LSH

- LSH for Jaccard distance
- LSH for Euclidean distance
- LSH for Hamming distance

# Summary

- Recommendation Systems
- Main Approaches
  - Content-Based Recommendations
  - Collaborative Filtering
- Approximate Nearest Neighbor Search
  - Locality Sensitive Hashing